

## ## 🔮 CATEGORY 1: APPLIED DEVOPS SCENARIOS

### ### Scenario 1: Direct Production Edits (nano/SSH)

**\*\*Problem\*\*:** Developers edit production directly using ` nano` /SSH

**\*\*Why Bad\*\*:**

- ❌ No version control → can't track/rollback changes
- ❌ No testing → bugs go to production
- ❌ No code review → security risks
- ❌ No audit trail → can't track who changed what
- ❌ Direct server access → security vulnerability

**\*\*First Steps\*\*:**

1. **\*\*Immediate\*\*:** Set up Git repository, move all code to version control
2. **\*\*CI/CD Pipeline\*\*:** Create CI/CD pipeline with stages: lint → test → build → deploy
3. **\*\*Infrastructure as Code\*\*:** Use Kubernetes manifests or Helm charts instead of manual edits
4. **\*\*Access Control\*\*:** Remove SSH access, require all changes via Git
5. **\*\*Monitoring\*\*:** Add health checks (` /healthz` , ` /readyz` ) and metrics (Prometheus) to detect issues early

**\*\*Key Points\*\*:**

- CI/CD pipeline should have: lint → test → security → build → container-build → deploy
- Use Helm for templating (separates logic from configuration)
- Health probes in Kubernetes deployments
- Container registry for images

---

### ### Scenario 2: Application Crashes on Traffic Spikes

**\*\*Problem\*\*:** App can't handle traffic spikes, needs autoscaling

**\*\*Solution - HPA (Horizontal Pod Autoscaler)\*\*:**

- **Configuration**: Min 2, Max 5 replicas, Target CPU 30%
- **How It Works**: Kubernetes monitors CPU, scales up when >30%, down when <30%
- **Metrics**: Uses Prometheus metrics (`/metrics` endpoint, port 9090)

**\*\*Additional Steps\*\*:**

1. **Health Probes**: Liveness (`/healthz`) and Readiness (`/readyz`) probes
2. **Stateless Design**: App must be stateless (use external database/storage, not local filesystem)
3. **Load Balancing**: Kubernetes Service distributes traffic across replicas
4. **Monitoring**: Grafana dashboard to visualize scaling

**\*\*Key Points\*\*:**

- HPA enabled in production with pod anti-affinity for fault tolerance
- Metrics exposed on port 9090 with Prometheus annotations
- Application must be stateless to scale horizontally

---

### ### Scenario 3: Secrets in Git Repository

**\*\*Problem\*\*:** Secrets hardcoded in Git, security risk

**\*\*Solution - Secrets Management\*\*:**

1. **Remove from Git**: Delete all secrets from codebase
2. **CI/CD Variables**: Store secrets in CI/CD variables (encrypted, masked)
3. **Injection at Deploy**: Use templating (Helm `--set-string`) to inject secrets at deploy time
4. **Kubernetes Secrets**: Secrets created in K8s, never in Git
5. **Rotation**: Easy to rotate (update CI/CD variable, redeploy)

#### **Implementation Flow**:

- Secrets in CI/CD variables: `DATABASE\_URL`, `JWT\_SECRET`, `MINIO\_ACCESS\_KEY`, etc.
- Templating command: `--set-string`  
serverSecret.extra.DATABASE\_URL="\${DATABASE\_URL}"`
- Kubernetes Secret created, pods mount as env vars
- **Never in Git**: Values files have empty placeholders, secrets only in CI/CD

---

### ### Scenario 4: No Monitoring/Visibility

**Problem**: Can't see what's happening in production

#### **Solution - Observability**:

1. **Metrics**: Prometheus metrics (4 golden signals)
2. **Health Checks**: Liveness and Readiness probes
3. **Logging**: Application logs via `kubectl logs`
4. **Visualization**: Grafana dashboards

#### **Implementation**:

- **Metrics Service**: Expose Prometheus metrics endpoint

- **Metrics**: `http\_request\_duration\_seconds`, `http\_requests\_total`, `http\_errors\_total`
- **Health**: `/healthz` (liveness), `/readyz` (readiness, checks DB)
- **Prometheus**: Port 9090, Service annotations `prometheus.io/scrape: "true"``
- **Grafana**: Dashboards with panels (latency, errors, CPU, pod health, HPA)

---

### ### Scenario 5: Manual Deployments

**Problem**: Deployments done manually, error-prone

**Solution - CI/CD Pipeline**:

1. **Automate**: CI/CD pipeline (e.g., GitLab CI/CD)
2. **Stages**: Lint → Test → Security → Build → Container → Scan → Deploy
3. **Atomic Deployments**: `--atomic` flag for automatic rollback
4. **Zero-Downtime**: Rolling updates (new pods ready before old removed)

**Pipeline Stages**:

- **Lint**: Code quality checks
- **Test**: Unit tests with coverage requirements
- **Security**: Secret detection, vulnerability scanning
- **Build**: Application builds
- **Container Build**: Docker images to registry
- **Image Scan**: Security scanning (Trivy)
- **Deploy**: Helm to Kubernetes, atomic with rollback

---

## ## 💡 CATEGORY 2: DEVOPS CONCEPTS

### ### Concept 1: Cattle vs. Pets

#### \*\*Definition\*\*:

- \*\*Pets\*\*: Named servers, manually maintained, irreplaceable (if dies = crisis)
- \*\*Cattle\*\*: Numbered servers, identical, easily replaced (if dies = replace it)

#### \*\*How It Influences Autoscaling\*\*:

- \*\*Stateless Design\*\*: Apps must be stateless (external database/storage, not local filesystem)
- \*\*Immutable Infrastructure\*\*: Containers built once, never modified (Docker images in CI/CD)
- \*\*Horizontal Scaling\*\*: Add/remove identical pods (HPA: 2-5 replicas)
- \*\*No Manual Intervention\*\*: Pods created/destroyed automatically (Kubernetes)
- \*\*Fault Tolerance\*\*: Pod crashes → Kubernetes creates new one (liveness probe)

#### \*\*Key Points\*\*:

- Stateless applications enable horizontal scaling
- Identical pods can be replaced without data loss
- HPA automatically creates/destroys pods based on metrics
- Pod anti-affinity ensures pods spread across nodes

---

### ### Concept 2: Idempotency

\*\*Definition\*\*: Operation produces same result whether executed once or multiple times.

#### \*\*Why Important\*\*:

- **Safe Retries**: Deployment fails → retry without side effects
- **Consistency**: Same config applied multiple times = same result
- **CI/CD Reliability**: Pipeline re-run without breaking things

#### **Examples**:

- `helm upgrade --install` (idempotent - can run multiple times)
- `kubectl apply` (idempotent - creates if missing, updates if exists)
- Database migrations (idempotent - safe to run multiple times)
- Docker builds (same image = same result)

#### **Non-Idempotent (Bad)**:

- Deleting file (can't delete twice)
- Incrementing counter (different result each time)

---

### ## Concept 3: GitOps

**Definition**: Infrastructure and code in Git, Git = single source of truth. Changes via Git commits, not manual commands.

#### **Principles**:

1. **Declarative**: Infrastructure in YAML (Kubernetes manifests, Helm charts)
2. **Version Controlled**: All config in Git (e.g., `k8s/`, `helm/` directories)
3. **Automated**: CI/CD applies changes (e.g., ` .gitlab-ci.yml` deploys on `main` )
4. **Observable**: Monitor what's deployed (Grafana dashboard)

#### **Implementation**:

- **Infrastructure as Code**: All K8s resources in version control

- **CI/CD Pipeline**: Auto-deploys on push to main branch
- **No Manual Changes**: All deployments via Git commits, not manual commands
- **Rollback**: Git history allows reverting to previous versions
- **Secrets**: Injected from CI/CD variables, not in Git

**Benefits:**

- Audit trail (who changed what, when)
- Easy rollback (revert Git commit)
- Consistency (same config in staging/production)

---

### ### Concept 4: Shift Left

**Definition**: Move testing, security, quality checks earlier in development lifecycle (left = earlier in timeline).

**Traditional (Bad)**: Test and fix issues in production

**Shift Left (Good)**: Test and fix issues during development

**CI/CD Pipeline Stages (Shift Left)**:

1. **Lint Stage** (earliest): Code quality checks, catches syntax errors early
2. **Test Stage**: Unit tests run before deployment, coverage requirements, failures block deployment
3. **Security Stage**: Scans for hardcoded secrets, catches security issues early
4. **Build Stage**: Builds fail fast if code doesn't compile, catches dependency issues early
5. **Image Scan Stage**: Scans Docker images for vulnerabilities
6. **Deploy Stage** (last): Only deploys if all previous stages pass

**\*\*Benefits\*\*:**

- Faster feedback (find issues in minutes, not days)
- Lower cost (fixing bugs in dev cheaper than production)
- Better quality (catch issues before users see them)

---

### ### Concept 5: Observability (Metrics/Logs/Traces)

**\*\*Definition\*\*:** Three pillars of observability

**\*\*Three Pillars\*\*:**

1. **Metrics**: Numerical measurements over time (CPU, memory, request rate)
2. **Logs**: Text records of events (application logs, error messages)
3. **Traces**: Request flow through distributed systems

**\*\*Implementation\*\*:**

**\*\*Metrics (Prometheus)\*\*:**

- **Four Golden Signals**:
- **Latency**: `http\_request\_duration\_seconds` (histogram, buckets 0.1s-10s)
- **Traffic**: `http\_requests\_total` (counter, labels: method, route, status)
- **Errors**: `http\_errors\_total` (counter, status >= 400)
- **Saturation**: Custom metrics (e.g., `posts\_total`, `users\_total`)
- **Exposure**: `/metrics` endpoint on port 9090
- **Scraping**: Prometheus scrapes via Service annotations (`prometheus.io/scrape: "true"`)

**\*\*Logs\*\*:**

- Application logs via `kubectl logs`
- Structured logging

**\*\*Health Checks\*\*:**

- **Liveness Probe**: `/healthz` (is app running?)
- **Readiness Probe**: `/readyz` (is app ready for traffic? checks database)

**\*\*Visualization\*\*:** Grafana dashboard with panels (latency, errors, CPU, pod health, etc.)

---

**### Concept 6: Immutable Infrastructure**

**\*\*Definition\*\*:** Servers/containers never modified after creation. To update, create new containers and destroy old ones.

**\*\*Traditional (Mutable)\*\*:** SSH into server, edit files, install packages → server drifts from original state

**\*\*Immutable\*\*:** Build new container image, deploy it, destroy old containers → always consistent

**\*\*How Containerization Enables It\*\*:**

- **Docker Images**: Built once in CI/CD (e.g., `build-server-image`, `build-client-image` jobs)
- **No Manual Changes**: Containers are read-only (except volumes)
- **Version Control**: Images tagged with Git commit (e.g., `IMAGE\_TAG: \$CI\_COMMIT\_REF\_SLUG`)
- **Rolling Updates**: Kubernetes replaces old pods with new ones (zero-downtime)

**\*\*Benefits\*\*:**

- Consistency (dev = staging = production)
- Reproducibility (same image = same behavior)
- Easy rollback (deploy previous image)
- Security (no manual changes = fewer vulnerabilities)

---

### ### Concept 7: CI/CD Pipelines

**Definition:** Automated pipeline that builds, tests, and deploys code.

**Pipeline Stages** (example):

1. **Lint Stage**: ESLint checks code, catches syntax/style errors early
2. **Test Stage**: Unit tests with coverage requirements, ensures code works before deployment
3. **Security Stage**: Scans for hardcoded secrets, prevents secrets in Git
4. **Build Stage**: Creates build artifacts, artifacts cached
5. **Container Build Stage**: Builds Docker images, pushes to registry, images tagged with commit SHA
6. **Image Scan Stage**: Scans images for HIGH/CRITICAL vulnerabilities
7. **Deploy Stage**: Deploys to Kubernetes using Helm, creates ConfigMaps, Secrets, Deployments, Services, Ingress

**Key Features**:

- **Atomic Deployments**: `--atomic` flag rolls back on failure
- **Secrets Injection**: Secrets from CI/CD variables, not in Git
- **Zero-Downtime**: Rolling updates (new pods ready before old ones removed)

---

### ### Concept 8: Version Control

**Definition:** System for tracking changes to files over time.

**Why Important:**

- **History:** See who changed what, when, why
- **Rollback:** Revert to previous versions
- **Collaboration:** Multiple developers work on same code
- **Branching:** Test changes without affecting main code

**Best Practices:**

- Commit frequently with clear messages
- Use branches for features
- Never commit secrets
- Use pull requests for code review

---

### ### Concept 9: Infrastructure as Code (IaC)

**Definition:** Managing infrastructure using code/configuration files instead of manual processes.

**Benefits:**

- **Version Control:** Infrastructure changes tracked in Git
- **Reproducibility:** Same config = same infrastructure
- **Automation:** CI/CD can deploy infrastructure
- **Consistency:** Dev = staging = production

**\*\*Tools\*\*:**

- **Kubernetes**: Declarative YAML manifests
- **Helm**: Templating engine for K8s (separates logic from config)
- **Docker Compose**: Local development orchestration

---

**### Concept 10: Security in DevOps**

**Definition**: Integrating security practices throughout the development lifecycle.

**Security Measures**:

1. **Secret Management**:

- Secrets in CI/CD variables (encrypted, masked)
- Never in Git
- Injected at deploy time via templating

2. **Security Scanning**:

- **Secret Detection**: Scans for hardcoded secrets in Git
- **Image Scanning**: Scans Docker images for vulnerabilities (HIGH/CRITICAL)

3. **Access Control**:

- RBAC in Kubernetes
- ServiceAccount for CI/CD
- No direct SSH access to containers

4. **Network Security**:

- NetworkPolicy for traffic control

- Ingress with TLS (cert-manager, Let's Encrypt)

## 5. \*\*Application Security\*\*:

- JWT authentication
- Password hashing (bcrypt)
- Input validation
- XSS prevention

**\*\*Shift Left Security\*\*:** Security checks early in pipeline (secret detection, image scanning)

---

## ## 🔎 CATEGORY 3: RETROSPECTIVE & ANALYSIS

> **\*\*Note\*\*:** This section contains project-specific details. Replace with your own project information.

### ### Your Project Summary

**\*\*Technical Stack\*\*:**

- **Frontend**: Next.js (React), TypeScript, Tailwind CSS
- **Backend**: NestJS (Node.js), TypeScript, Prisma ORM
- **Database**: PostgreSQL (Kubernetes deployment)
- **Storage**: MinIO (S3-compatible object storage)
- **Containerization**: Docker (multi-stage builds)
- **Orchestration**: Kubernetes (Deployments, Services, Ingress, PVCs)
- **Templating**: Helm (separates logic from configuration)

- **CI/CD**: GitLab CI/CD (7 stages: lint → test → security → build → container-build → image-scan → deploy)
- **Monitoring**: Prometheus (metrics), Grafana (dashboards)
- **Security**: Trivy (image scanning), gitleaks (secret detection)

#### **Architecture Decisions**:

- **Stateless Design**: No local filesystem storage (PostgreSQL + MinIO external)
- **Health Probes**: Liveness (`/healthz`) and Readiness (`/readyz`)
- **Autoscaling**: HPA (2-5 replicas, CPU-based, production only)
- **Fault Tolerance**: Pod anti-affinity (pods on different nodes)
- **Secrets Management**: GitLab CI variables → Helm → Kubernetes Secrets
- **Templating**: Helm (values.yaml for staging, values-prod.yaml for production)

---

### Question: "What technical stack did you use? What would you change?"

#### **Current Stack**:

- **Frontend**: Next.js (good for SSR, but adds complexity)
- **Backend**: NestJS (excellent for large apps, but might be overkill for simple CRUD)
- **Database**: PostgreSQL (perfect choice, relational data)
- **Storage**: MinIO (good S3-compatible, but AWS S3 might be better for production)
- **Orchestration**: Kubernetes + Helm (industry standard, but complex for small teams)
- **CI/CD**: GitLab CI (good, but could add more stages like integration tests)

#### **What I Would Change**:

1. **\*\*Simplify Frontend\*\*:** Consider React without Next.js if SSR not needed (faster builds)
2. **\*\*Add Integration Tests\*\*:** E2E tests in CI/CD pipeline (currently only unit tests)
3. **\*\*Use Managed Services\*\*:** Consider managed PostgreSQL (AWS RDS) instead of self-hosted
4. **\*\*Add Monitoring Alerts\*\*:** Set up Prometheus alerts (currently only dashboards)
5. **\*\*Improve Security\*\*:** Add OWASP dependency scanning (currently only Trivy for images)
6. **\*\*Optimize Images\*\*:** Use multi-stage builds more aggressively (reduce image size)
7. **\*\*Add Blue-Green Deployment\*\*:** For zero-downtime (currently only rolling updates)

**\*\*Why These Changes\*\*:**

- **\*\*Managed Services\*\*:** Less operational overhead, better reliability
- **\*\*More Testing\*\*:** Catch integration issues before production
- **\*\*Alerts\*\*:** Proactive issue detection (not just reactive dashboards)
- **\*\*Smaller Images\*\*:** Faster deployments, lower costs

---

### Question: "How did you ensure your application is stateless? Why is this important?"

**\*\*How We Ensured Statelessness\*\*:**

1. **\*\*Database Externalization\*\*:**

- All data in PostgreSQL (Kubernetes service, not local)
- ORM connects to service name (e.g., `postgresql:5432`, not localhost)
- Data persists in PersistentVolumeClaim (survives pod restarts)

## 2. \*\*File Storage Externalization\*\*:

- Images uploaded to MinIO (S3-compatible)
- Storage service uses `PutObject` / `GetObject` (never writes to `/tmp`)
- Files accessible from any pod (shared storage)

## 3. \*\*Session Management\*\*:

- JWT tokens stored client-side (browser localStorage)
- Sessions in database (not in-memory)
- No pod-local session storage

## 4. \*\*No Local State\*\*:

- No files written to container filesystem
- No in-memory caches (all data from database/storage)

### \*\*Why Important\*\*:

- \*\*Horizontal Scaling\*\*: Multiple pods can serve requests (no data inconsistency)
- \*\*Pod Failures\*\*: If pod crashes, another pod takes over (no data loss)
- \*\*Rolling Updates\*\*: Zero-downtime deployments (old pods removed after new ones ready)
- \*\*Chaos Engineering\*\*: Application survives random pod kills

### \*\*Verification\*\*:

- Pods can be deleted and recreated without data loss
- HPA scales to multiple replicas, all serve same data
- Rolling updates work without downtime

---

### Question: "Explain your secrets management strategy. Why is it secure?"

**Strategy:** Never commit secrets to Git. All secrets stored in CI/CD variables, injected only during deployment.

**Implementation:**

1. **Storage:**

- Secrets in CI/CD variables (encrypted at rest, masked in logs)
- Variables: `DATABASE\_URL`, `JWT\_SECRET`, `MINIO\_ACCESS\_KEY`, `MINIO\_SECRET\_KEY`, `NEXTAUTH\_SECRET`

2. **Injection:**

- Templating uses empty placeholders in values files (not in Git)
- CI/CD job passes secrets via `--set-string` flags
- Secrets only exist in: CI/CD variables → CI job → Kubernetes Secrets → Pods

3. **Security Measures:**

- **Masked:** Secrets masked in CI/CD logs (can't see values)
- **Protected:** Variables marked as "protected" (only available in protected branches)
- **Base64 Encoding:** Kubernetes Secrets base64-encoded (not plaintext in etcd)
- **No Git History:** Secrets never in Git, so no risk of exposure in history

**Why Secure:**

- **No Git Exposure:** Even if repo is public, secrets aren't exposed
- **Access Control:** Only CI/CD job can access secrets (not developers)
- **Audit Trail:** CI/CD logs who accessed secrets (if enabled)
- **Rotation:** Easy to rotate (update CI/CD variable, redeploy)

**\*\*Example Flow\*\*:**

1. Developer pushes code (no secrets in code)
2. CI/CD job reads `JWT\_SECRET` from CI/CD variables
3. Templating command: `--set-string  
serverSecret.extra.JWT\_SECRET="\${JWT\_SECRET}"`
4. Kubernetes Secret created (base64-encoded)
5. Pod mounts secret as environment variable

---

### Question: "How did you implement monitoring? What metrics do you track?"

**\*\*Implementation\*\*:**

1. **Metrics Service** (e.g., `server/src/metrics/metrics.service.ts`):
  - **Four Golden Signals**:
    - **Latency**: `http\_request\_duration\_seconds` (histogram, buckets: 0.1s-10s)
    - **Traffic**: `http\_requests\_total` (counter, labels: method, route, status)
    - **Errors**: `http\_errors\_total` (counter, status >= 400)
    - **Saturation**: Custom metrics (e.g., `posts\_total`, `users\_total`)
  - **Exposure**: Separate server on port 9090, `/metrics` endpoint
  - **Format**: Prometheus text format
2. **Metrics Middleware**:
  - Intercepts all HTTP requests
  - Records duration, method, route, status code
  - Calls metrics service

### 3. \*\*Kubernetes Integration\*\*:

- \*\*Service Annotations\*\*: `prometheus.io/scrape: "true"`, `prometheus.io/port: "9090"`
- \*\*NetworkPolicy\*\*: Allows Prometheus from monitoring namespace
- \*\*Scraping\*\*: Prometheus scrapes every 30 seconds

### 4. \*\*Health Checks\*\*:

- \*\*Liveness\*\*: `/healthz` (app running?)
- \*\*Readiness\*\*: `/readyz` (app ready for traffic? checks database)

### 5. \*\*Visualization\*\*:

- \*\*Grafana Dashboard\*\*: Multiple panels (latency, errors, CPU, pod health, HPA, etc.)
- \*\*Thresholds\*\*: Color-coded (green/yellow/orange/red)
- \*\*Time Range\*\*: Configurable (e.g., 6 hours default)

#### \*\*What We Track\*\*:

- Request latency (95th percentile)
- Request rate (requests/second)
- Error rate (errors/second)
- Pod health (ready/unready)
- CPU usage (for HPA)
- Database status
- Storage usage

#### \*\*Use Cases\*\*:

- \*\*Performance\*\*: Identify slow endpoints (latency > 2s)
- \*\*Reliability\*\*: Detect errors (error rate > 0.1/s)

- **Scaling**: HPA uses CPU metrics (scale when > 30%)
- **Debugging**: Pod restarts indicate crashes

---

### Question: "Why did you choose Helm over Kustomize? (or vice versa)"

**Helm vs. Kustomize**:

**Helm**:

- **Pros**:
  - Templating engine (Go templates)
  - Values files (separates logic from config)
  - Package manager (can share charts)
  - Atomic deployments (`--atomic` flag)
  - Industry standard, large community
- **Cons**:
  - More complex (templates can be hard to read)
  - Requires Helm binary
  - Steeper learning curve

**Kustomize**:

- **Pros**:
  - Native to Kubernetes (built into `kubectl`)
  - Simpler (YAML overlays)
  - No external tools needed
  - Easy to understand
- **Cons**:

- Less powerful (no templating)
- Harder to share/reuse
- No atomic deployments

#### **\*\*Why I Chose Helm\*\*:**

- **\*\*Templating\*\*:** Need to separate logic from config (values.yaml vs templates)
- **\*\*Multiple Environments\*\*:** Staging and production (values.yaml vs values-prod.yaml)
- **\*\*Secrets Injection\*\*:** Easy to inject secrets via `--set-string`
- **\*\*Atomic Deployments\*\*:** `--atomic` flag for automatic rollback
- **\*\*Industry Standard\*\*:** More widely used, better documentation